

## 1. Transitioning from Regression to Classification

- Covers how to identify when a problem should be modeled using **classification** (predicting discrete labels) instead of **regression** (predicting continuous values).

## 2. Key Differences

- **Regression** predicts numeric values (e.g., house prices).
- **Classification** predicts discrete categories (e.g., email is spam or not).
- Introduces **binary classification** (two classes) and **multi-class classification** (more than two classes).

## 3. Data Preparation

- Shows how to convert a continuous target into **categorical classes**, such as binning age into groups or profit margins into “low/medium/high”.

## 4. Choosing the Right Model

- Presents common classification algorithms:
  - **Logistic Regression**
  - **K-Nearest Neighbors (KNN)**
  - **Decision Trees**
  - **Support Vector Machines (SVM)**
- Highlights that regression models unlike logistic models aren’t suitable for discrete class outputs.

## 5. Training & Model Evaluation

- Foreshadows splitting data into **train/test sets**, fitting classification models, and evaluating performance

## 1. Introduction to the ML Lifecycle

- Explains the purpose: helps developers **navigate the complexities** of building ML solutions.

## 2. Key Stages of the Lifecycle

### A. Understanding ML & Gathering Insights

- Defines ML as the ability for systems to **learn without explicit programming**.
- Highlights importance of grasping the foundational concept: learning from data.

### B. Framing the Business Problem

- Emphasizes beginning with the right question.
- Must define business metrics and objectives *before* picking a model.

### C. Data Collection & Understanding

- Gather relevant data.
- Use exploratory data analysis (EDA) to understand structure, patterns, and quality.

### D. Feature Engineering (Preparation)

- Manipulate raw data to improve model input:
  - Create new meaningful features

- Handle correlations and missing values
- Encode categorical data through techniques like one-hot encoding (e.g., transforming 'ocean proximity' into binary flags)

## E. Building the ML Pipeline

- Combine preprocessing, feature engineering, and model training into a **reproducible pipeline**.
- Helps maintain consistent transformations throughout experimentation and deployment.

## F. Model Training & Evaluation

- Train models and evaluate them on key metrics:
  - Classification: accuracy, precision, recall, confusion matrix
  - Regression: RMSE, MAE
- Compare against baseline models.

## G. Operationalizing the Pipeline

- Deploy the end-to-end pipeline into production:
  - Ensure it's **scalable, maintainable**, and version-controlled.
  - Include monitoring for data drift and retraining.